

Zero to Cloud-Native with IBM Cloud

Part 9A: Creating a Tekton Pipeline

Kevin Collins

Technical Sales Leader

IBM Cloud Enterprise Containers – Americas

Kunal Malhotra

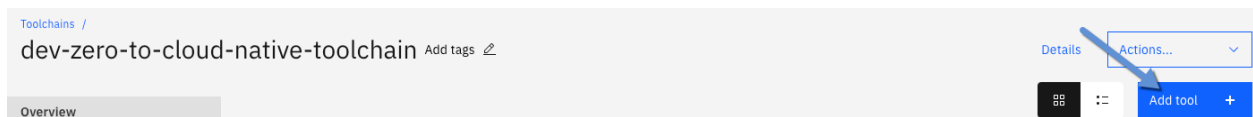
Cloud Platform Engineer

IBM Cloud MEA

1 – IBM Cloud Continuous Delivery – Classic Pipeline

Note: This tutorial will assume you have already setup a toolchain as describe in Part 9 – Section 4.

This tutorial will go step through creating is a classic pipeline for the API frontend microservice. To get started, navigate to the zero-to-cloud-native-toolchain you created in part 9 for this series.



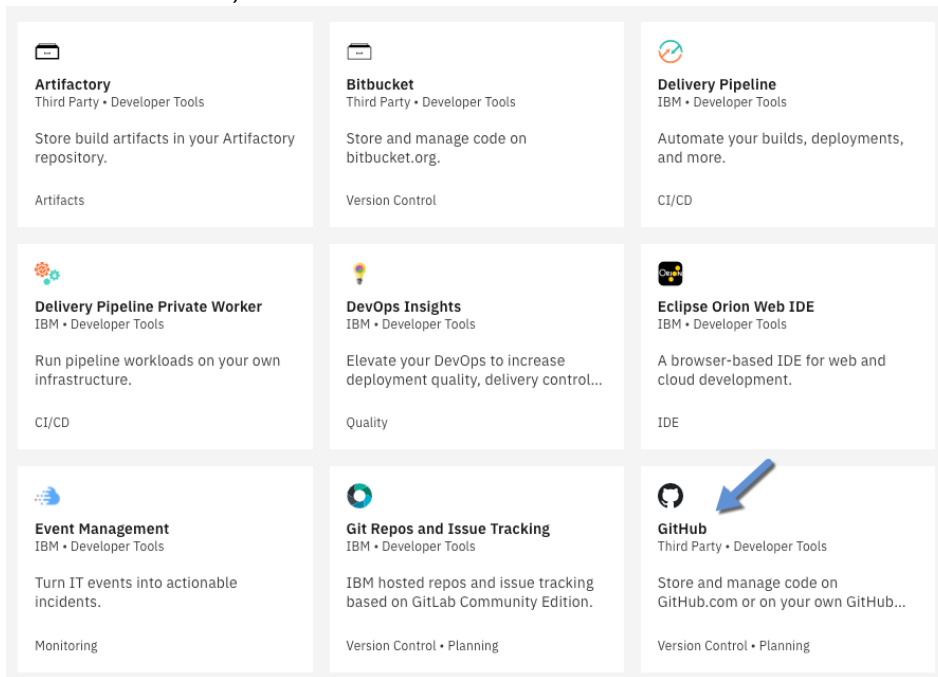
1 – 1 GitHub Integration

The first thing you need to do is connect your SCM (Source Code Management) repository with the toolchain. You do this by clicking add tool and then selecting SCM (GitHub, Bitbucket, GitLab, or GitHub Enterprise Whitewater).

Note: Repositories can be Public or Private. In this tutorial, I will be the public repository that I've shared and that you clone in part 8. By default, I recommend setting your repositories to private unless you plan to share them publicly. This way, if you make a mistake and upload something like an API key, you are still protected.

Note: if this is the first time you are creating a toolchain with a GitHub repository, you will first need to Authorize IBM Cloud to access GitHub. You can follow the instructions here if needed: <https://cloud.ibm.com/docs/ContinuousDelivery?topic=ContinuousDelivery-integrations#github>

On the next screen, select **GitHub**.



Select **Existing** as the repository type and enter the repository URL for your git repository that you setup in Part 8 of this tutorial series. Click **Create Integration** after entering these settings.



The image shows a form for creating a GitHub integration. The form has several sections:

- GitHub Server:** A dropdown menu showing "GitHub (https://github.com)".
- Authorized as kmcoll with access granted to zero GitHub organization(s)** with a link to [Manage Authorization](#).
- Repository type:** A dropdown menu showing "Existing".
- Link to the repository that is specified in the Repository URL field.**
- Repository URL:** A text input field containing "https://github.com/kmcoll/api-frontend-02cn".
- Integration Owner:** A dropdown menu showing "kmcoll".
- Enable GitHub Issues:** A checked checkbox.
- Track deployment of code changes:** An unchecked checkbox.

1 – 2 Create Delivery Pipeline

The next tool we need to add is a Delivery Pipeline. Click on **add tool** and then click on the **Delivery Pipeline** tile.



Delivery Pipeline

IBM • Developer Tools

Automate your builds, deployments, and more.

CI/CD

On the next screen, give your delivery pipeline a name. I will use **dev-classic-frontend-api-toolchain**. For this part of the tutorial, keep **Classic** selected as the Pipeline type, and click **Create Integration**.

[Integration](#) /

Pipeline



Pipeline name:


Pipeline type: Classic

Classic pipelines provide an easy to use graphical UI for defining stages and jobs that run on public shared infrastructure, with support for running individual stages on Private Workers.

☒ Show apps in the View app menu ⓘ

1 – 2 – 1 Add the Continuous Delivery Service

On the next screen, you will see all the tools we have added to our toolchain. The next step is to configure the delivery pipeline we just created. You will also see the following warning notice.

 **Continuous Delivery service required:**
A service instance was not found in this resource group and region us-south. If you do not take action before 2020-09-07, 14:02 UTC, Delivery Pipeline jobs will not run, pushes to Git Repos and Issue Tracking will fail, and DevOps Insights will be disabled. [Add the service](#) to the resource group to ensure uninterrupted use of the service. [Learn more](#). For more information about usage and billing, see the [blog](#).

You have 5 days to use the service for free. If you want to use the service more than 5 days, click on **Add the service**.

The first step on the next page is to select a region and plan. Keep the region as **Dallas**, like all our other resources, and you can also keep the plan selected as **Lite**.

Catalog / Services /

Continuous Delivery

Author: IBM • Date of last update: 05/01/2020 • Docs

Create About

Select a region

Select a region

Dallas

Select a pricing plan

Displayed prices do not include tax. Monthly prices shown are for country or region: United States

Plan	Features	Pricing
Lite	Continuous Delivery for organizations (orgs) or resource groups of up to 5 users 5 users per organization or resource group 500 Delivery Pipeline jobs run per organization per month or per resource group per month 500 MB of private Git Repos storage per organization or resource group The Lite plan offers the full capabilities of Continuous Delivery, with usage limits, to small teams at no cost. Lite plan services are deleted after 30 days of inactivity.	Free

Scrolling down, you need to give you service a name and indicate which resource group to place it in. I will use **Continuous Delivery-zero-to-cloud-native** as the name and select **zero-to-cloud-native** as the resource group.

After entering those settings, click on **Create**.

Configure your resource

Service name

Continuous Delivery-zero-to-cloud-native

Select a resource group ⓘ

zero-to-cloud-native

Tags ⓘ

Examples: env:dev, version-1

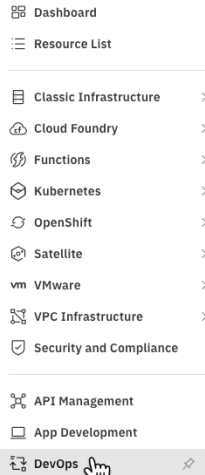
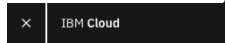
Create

Add to estimate

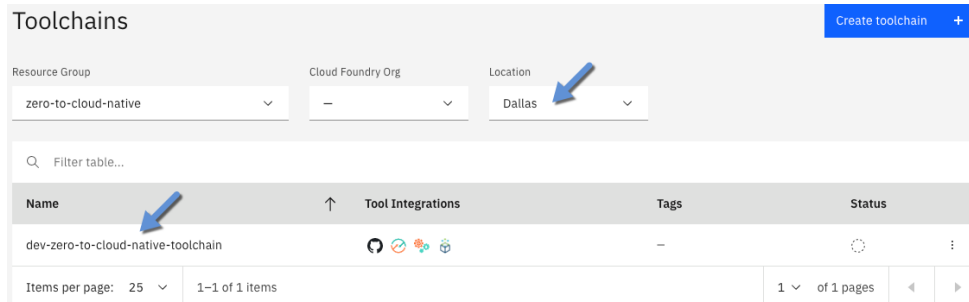
View terms

1 – 3 Configure your Delivery Pipeline

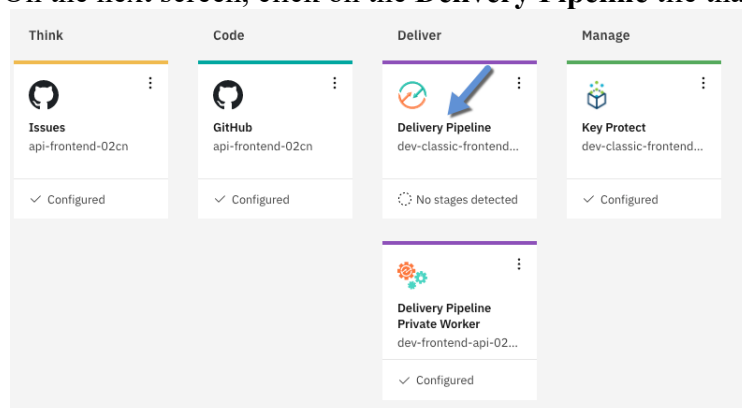
Navigate back to the toolchain you just created by clicking on the IBM Cloud hamburger menu and **select DevOps**.



On the next screen, select **Dallas** as the location your toolchain is in and then click on the **toolchain name**.

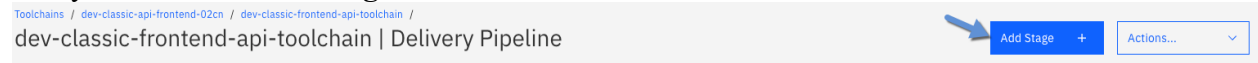


On the next screen, click on the **Delivery Pipeline** tile that you just created.



1 – 3 - 1 Build Stage

This will bring you to what looks like an empty screen as your pipeline does not have any stages in it yet. Click on **Add Stage**.



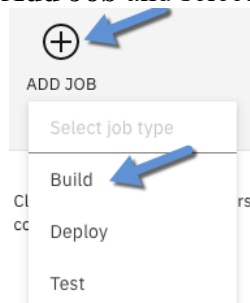
Start by naming the first stage **Build**.

Your git hub repository should already be populated as we have already created Git integration. Next, we will be building a dev environment, select the **dev branch** that you created in part 8 of this tutorial series.

Make sure to keep the '**When a commit is pushed**' option selected. This will automatically start our toolchain once a commit is pushed to GitHub.

A screenshot of the 'BUILD' stage configuration form. At the top, there is a 'Delete' button. Below it are tabs for 'Input', 'Jobs', 'Workers' (with a 'New' badge), and 'Environment pro...'. The 'Input' tab is active. Under 'Input settings', there is an 'Input type' dropdown set to 'Git repository'. Below that is a 'Git repository' dropdown set to 'api-frontend-02cn'. The 'Git URL' field contains 'https://github.com/kmcoll/api-frontend-02cn.git'. The 'Branch' dropdown is set to 'dev'. Under 'Stage trigger', the radio button 'Run jobs automatically for Git events on the chosen branch' is selected, and the checkbox 'When a commit is pushed' is checked. At the bottom, there is an unchecked checkbox 'Allow this stage to be run manually by all toolchain members'. 'Cancel' and 'Save' buttons are at the bottom left. Blue arrows point to the 'Jobs' tab, the 'Git repository' dropdown, the 'When a commit is pushed' checkbox, and the 'Save' button.

Don't click Save quite yet, instead click on the **Jobs** tab. A stage in a classic toolchain consists of a number of jobs. The first job we will create is fetching the code from GitHub. Click on **Add Job** and select **Build**.



Change the name of the Job to **Fetch Code**, the builder type to **shell script** and enter

```
source ./pipeline/scripts/fetch_code.sh
```

as the Build script.

BUILD Delete

Input **Jobs** Workers New Environment pro...

Fetch Code + ADD JOB

Fetch Code Remove

Build configuration

Builder type ?
Shell Script ▼

Pipeline image version ?
Inherited from Configure Pipeline (latest) ▼

Build script ?
source ./pipeline/scripts/fetch_code.sh

Note: you can view all the scripts used in this tutorial from the **pipeline/scripts** directory in each git repository:

After entering these settings, click on **Add Job** and **Test** and name the job **Unit Test**. This job will make sure your code repository was successfully cloned.

Keep the job as **simple** and enter, set the builder type to **shell script** and enter

```
source ./pipeline/scripts/unit_test.sh
```

for Test script.

BUILD Delete

Input **Jobs** Workers New Environment pro...

Fetch Code > **Unit Test** + ADD JOB

Unit Test Remove

Test configuration

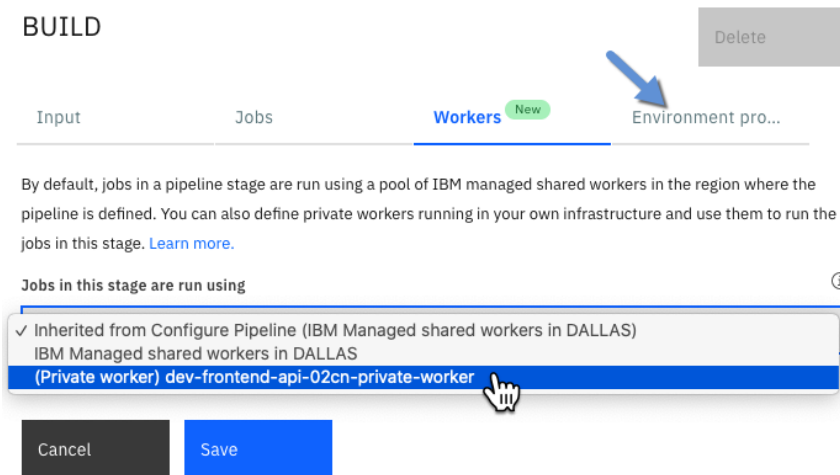
Tester type ?
Simple ▼

Pipeline image version ?
Inherited from Configure Pipeline (latest) ▼

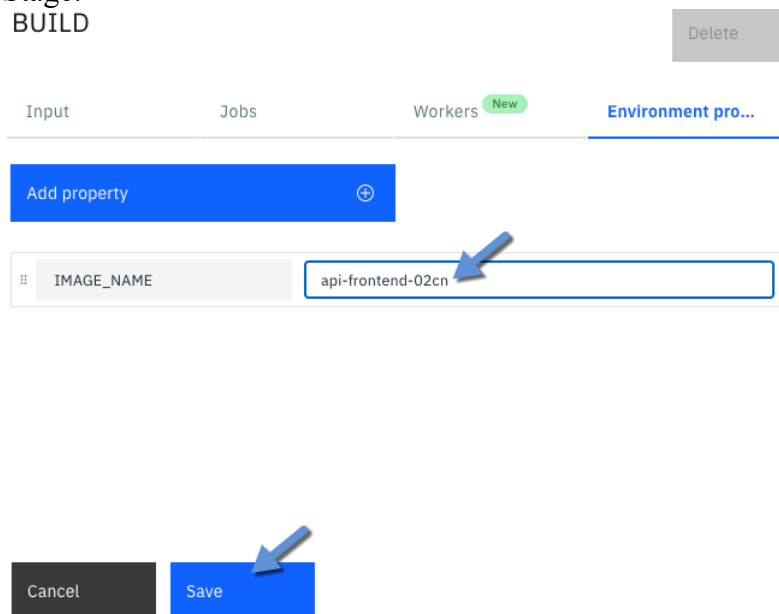
Test script ?
source ./pipeline/scripts/unit_test.sh

Best Practice: You will note that we are using scripts stored in our git repository to perform the various tasks for stage jobs versus entering the script code directly in the Job. It is much easier to maintain and develop these scripts in a single place outside of the toolchain itself. Furthermore, once we commit changes to our scripts, the commit will invoke our pipeline with the new scripts.

To configure your pipeline with the private worker you created, click on the **Workers tab** and select the private worker you previously created.



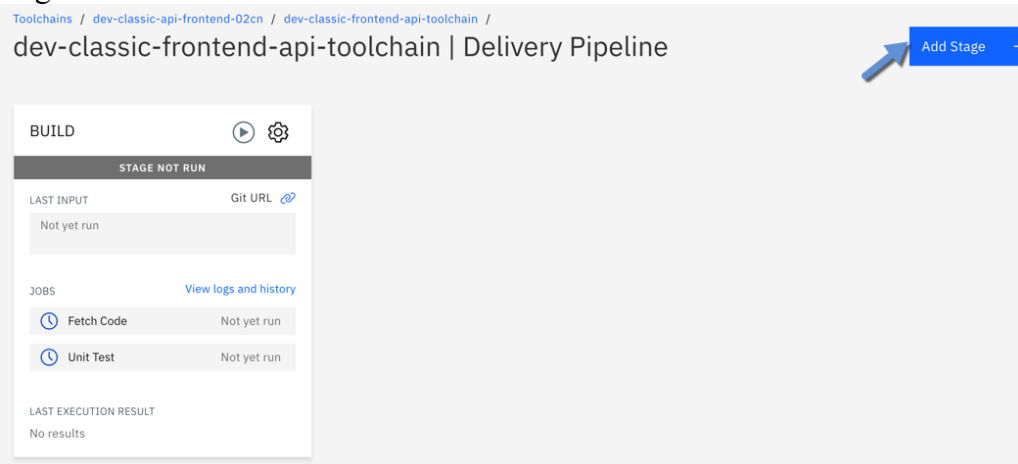
After selecting your private worker, the last step for the Build stage is to set an environment variable. Click on the **Environment properties** tab, click **Add Property** and select **Text Property**. Enter **IMAGE_NAME** as the property name and the name of the image for your code to be deployed as the value. For this example, we are building a CI/CD pipeline for the api frontend, so I will enter **api-frontend-02cn** as for image name. Click **Save** to save the Build Stage.



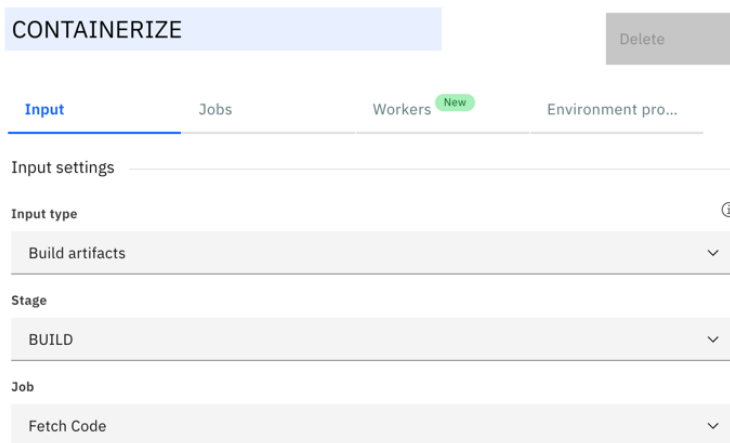
That completes the build stage.

1 – 3 – 2 Containerize Stage

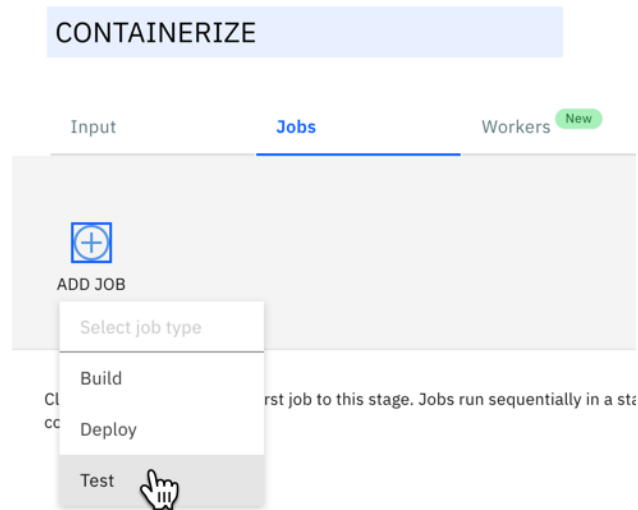
The next stage we will create is the containerization stage. Click on **Add Stage** to create a new stage.



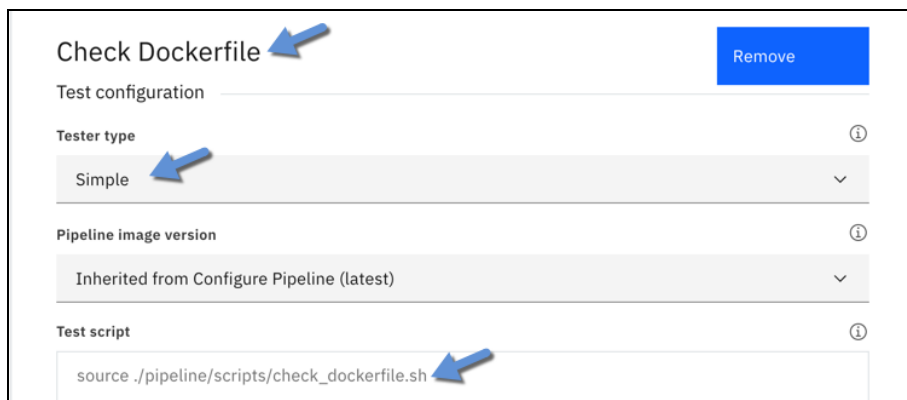
Change the name to **CONTAINERIZE**, leave the rest of the Input step as-is. On the **input tab**, **Build artifacts** should be the input type, **BUILD** as the stage and **Fetch Code** for the job. This will pass the code that was cloned from the git branch so we can create an image of the code in this step.



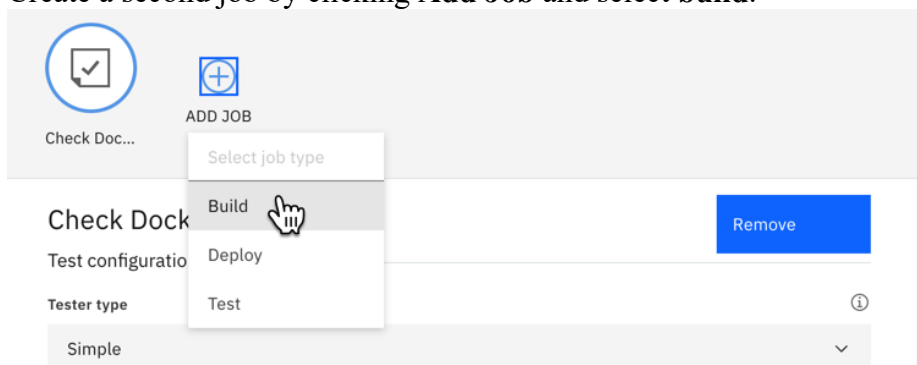
Next, we need to create the jobs that will build our container. The first job is making sure the Dockerfile is formatted correctly. Click on **Add Job** and select **Test**. Click on the **Jobs** tab then click **Add Job** and select **Test**.



Change the name to **Check Dockerfile**.
Keep the tester type as Simple and for Test script enter:
`source ./pipeline/scripts/check_dockerfile.sh`



Create a second job by clicking **Add Job** and select **build**.



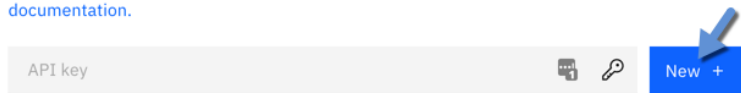
Name the job **Check Container Registry**. This job will make sure you have the appropriate permissions and quota to add new images to your IBM Cloud Registry.

Enter **Container Registry** of the Builder Type. This job requires an IBM Cloud API Key in order to access the **IBM Cloud Container Registry Service**. Input your **IBM Cloud API Key** or we will be following best practices and use **Key Protect** to store our API Key. Since this is the first time we will be using an IBM Cloud API Key in the toolchain service we will create a new API key which we will store in Key Protect. Click the **New** icon.

Enter an API key

This API key will be used to access the container registry and its namespaces.

Note: To maintain the strongest security posture you should use a service ID API key that has been configured to have only the access you need. For more information, see the [IAM documentation](#).



On the next screen, give the key a meaningful name. I will name the key **API Key for zero to cloud native toolchain**.

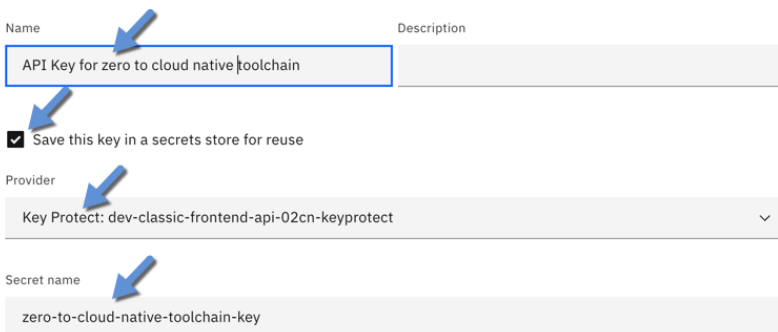
Check the box **Save this key in secrets for reuse**.

Under Provider, select the **Key Protect** instance we created in Part 3.

Finally, provide a **Secret Name** for the secret that will be stored in Key Protect. I will use **dev-classic-api-frontend-02cn** as my secret name.

Create a new API key with full access

anything you could do. You can improve your security posture by using the [IAM UI to create a service ID API key](#) that limits access to only what your pipeline requires, and then pasting that into the template UI instead. For more information on API keys and access see the [IAM documentation](#).



After entering all of these settings click OK.




On the next screen, click Authenticate.

Enter an API key

This API key will be used to access the container registry and its namespaces.


Note: To maintain the strongest security posture you should use a service ID API key that has been configured to have only the access you need. For more information, see the [IAM documentation](#).

.....



New +

Cancel

Authenticate 

Change the **IBM Cloud Region** to US South which is where our container registry is deployed. Select the **container registry namespace** you created earlier. The container registry namespace I created was zero-to-cloud-native. For the **docker image**, enter **api-frontend-02cn** since this pipeline is building the frontend api microservice.

Finally, for the build script, enter:

```
source ./pipeline/scripts/unit_test.sh
```

Container Registry

Remove

Build configuration

Builder type

Container Registry

Pipeline image version

Inherited from Configure Pipeline (latest)

API key

API Key for zero to cloud native toolchain

IBM Cloud region

US South

Account Name

kevin collins's Account

Container Registry namespace

zero-to-cloud-native

Docker image name

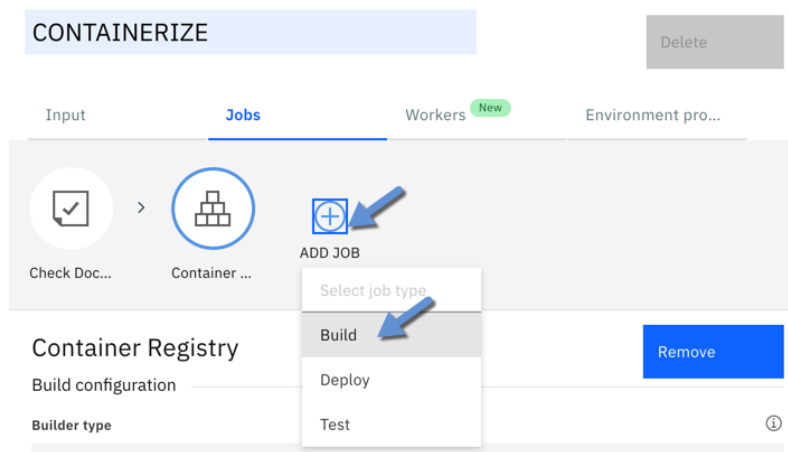
api-frontend-02cn

Build script

source ./pipeline/scripts/unit_test.sh

Don't click Save quite yet, we need to create one more job for this stage.

We need to create one more job to create the container image. To do so, click on **Add Job** and select **Build**.



Enter **Build Container Image** as the name and **Container Registry** as the **Builder type**. With the previous job we added an API key for our toolchain service to our Key Protect instance which we will be able to use. Click on the **api key** box, and then this time click on the **key icon**.

Enter an API key

This API key will be used to access the container registry and its namespaces.

Note: To maintain the strongest security posture you should use a service ID API key that has been configured to have only the access you need. For more information, see the [IAM documentation](#).



On the next screen, select the **Key Protect instance** you created and the **Secret Name** we create in the previous step. Then click **OK**.

API key

Select a secret for this field from a secrets store using the options below.

Provider

Key Protect: dev-classic-frontend-api-02cn-keyprotect

Secret name

zero-to-cloud-native-toolchain-key




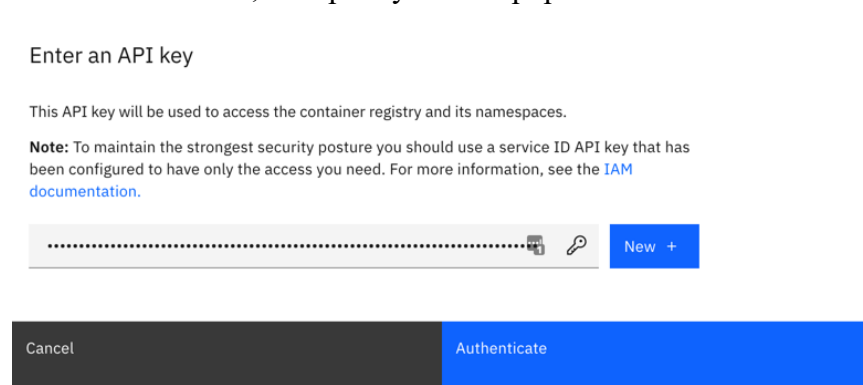
On the next screen, the api key will be populated. Click **Authenticate**.

Enter an API key

This API key will be used to access the container registry and its namespaces.

Note: To maintain the strongest security posture you should use a service ID API key that has been configured to have only the access you need. For more information, see the [IAM documentation](#).

.....



New +

Cancel

Authenticate

Just like we did on the previous job, you will need to change the **IBM Cloud Region** to US South which is where our container registry is deployed. Select the **container registry namespace** you created earlier. The container registry namespace I created was zero-to-cloud-native. For the **docker image**, enter **api-frontend-02cn** since this pipeline is building the frontend api microservice. For the build script, enter:

```
source ./pipeline/scripts/build_container.sh
```

Build Container Image

Remove

Build configuration

Builder type

Container Registry

Pipeline image version

Inherited from Configure Pipeline (latest)

API key

API Key for zero to cloud native toolchain

IBM Cloud region

US South

Account Name

kevin collins's Account

Container Registry namespace

zero-to-cloud-native

Docker image name

api-frontend-02cn

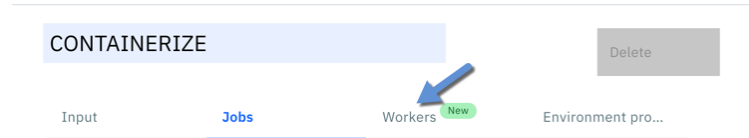
Build script

source ./pipeline/scripts/build_container.sh

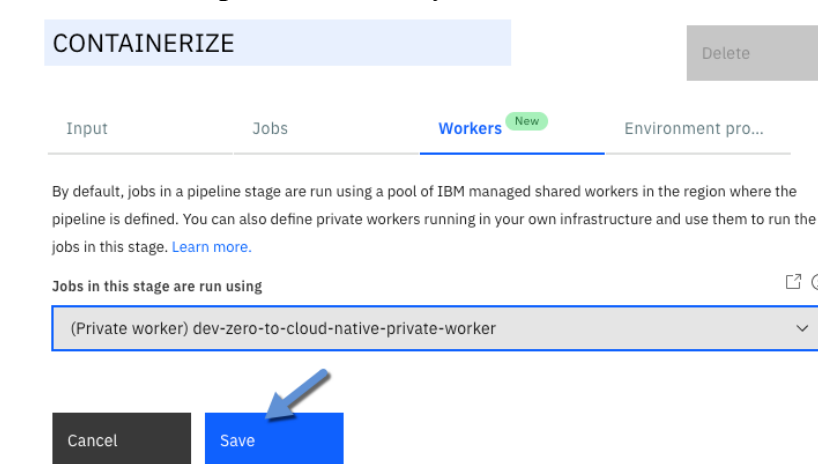
This finishes setting up the jobs, but there are a couple others things we need to do before saving the stage.

The next thing we need to do is to configure the containerize stage to use the private worker we created.

Click on the **workers** tab:

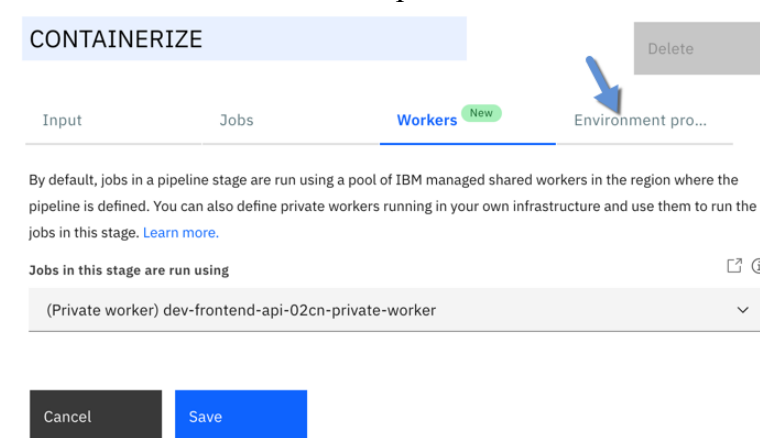


Next, select the **private worker** you created.

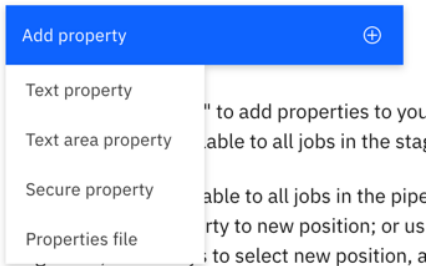


Before clicking save, we need to first set a couple environment properties.

Click on the Environment Properties tab.



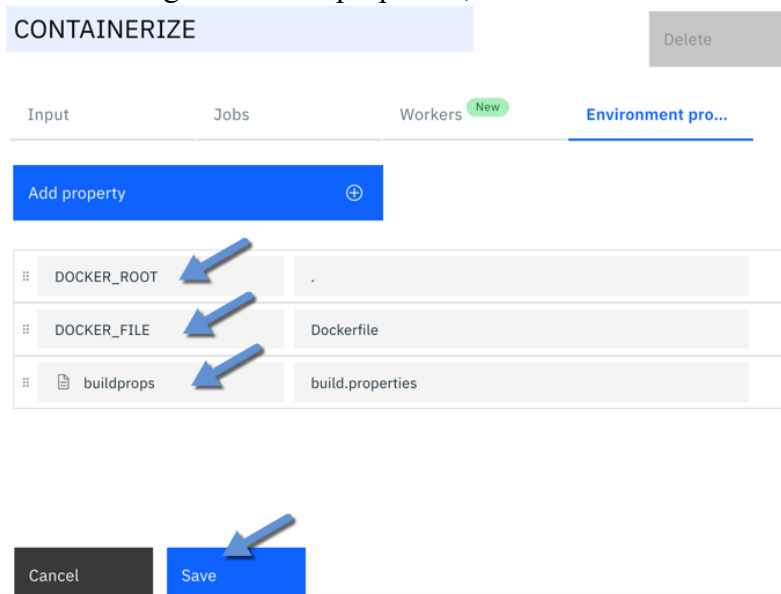
There are three Environment Properties we will need to add. To add a property, click on Add Property and enter the property type.



Referring to the table below, create the following environment properties.

Property Type	Property Name	Value
Text Property	DOCKER_FILE	Dockerfile
Text Property	DOCKER_ROOT	.
Properties File	buildprops	build.properties

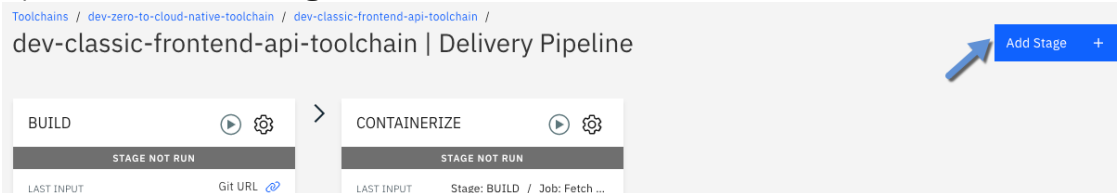
After entering all of these properties, click Save.



This completes the containerize job.

1 – 3 – 3 Deploy Stage

The final stage we need to create is the **deploy stage** where we will deploy the container into our OpenShift. Click **Add Stage**.

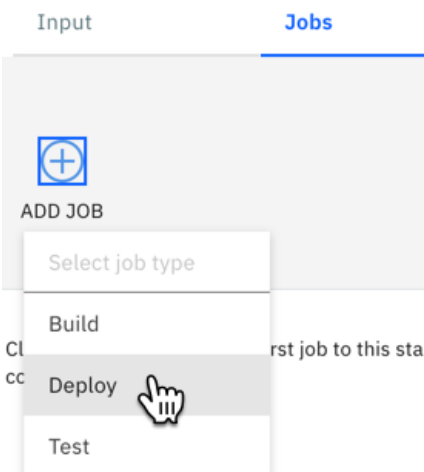


On the next screen name the stage **DEPLOY**. While on the **Input tab**, select **CONTAINERIZE** as the Stage to get inputs from and select the job **Build Container Image**. This will pass the container image that we will deploy to the deploy stage.



Next we need to add two more jobs. The first one to deploy the image and the second to make sure the deployment was successful. To create the Deploy to OpenShift Job, click on **Add Job**, select **Deploy**.

DEPLOY

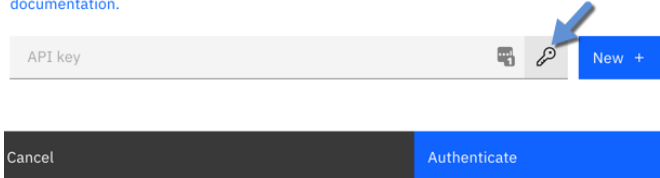


On the next screen, name the job **Deploy to OpenShift**. Select **Kubernetes** as the Deployer. Just like we did on the **Containerize Job**, click on the API Key text box, and then click on the Key Protect key icon.

Enter an API key

This API key will be used to access the cluster where your application will be deployed.

Note: To maintain the strongest security posture you should use a service ID API key that has been configured to have only the access you need. For more information, see the [IAM documentation](#).



On the next screen, enter your **Key Protect instance** and **Secret Name** for the toolchain service and click **OK**.

API key

Select a secret for this field from a secrets store using the options below.



The next screen will populate the API Key automatically, click **Authenticate**.

Enter an API key

This API key will be used to access the cluster where your application will be deployed.

Note: To maintain the strongest security posture you should use a service ID API key that has been configured to have only the access you need. For more information, see the [IAM documentation](#).



Scrolling down, enter US South as the **IBM Cloud Region**. Select the **Resource Group** where your cluster is created which should be **zero-to-cloud-native**. Next, selected your **cluster name** which should be **zero-to-cloud-native**.

Finally, for the deploy script, enter:

```
source ./pipeline/scripts/deploy_to_kubernetes.sh
```

Deploy to OpenShift Remove

Deploy configuration

Deployer type ⓘ
Kubernetes

Pipeline image version ⓘ
Inherited from Configure Pipeline (latest)

API key ⓘ
API Key for zero to cloud native toolchain

IBM Cloud region ⓘ
US South

Account Name ⓘ
kevin collins's Account

Resource Group ⓘ
zero-to-cloud-native

Cluster name ⓘ
zero-to-cloud-native

Deploy script ⓘ
source ./pipeline/scripts/deploy_to_kubernetes.sh

The last step in our pipeline is to validate that the deployment was successful. To do so, scroll back to the top of the page and click **Create Job** and select **Deploy**.

DEPLOY

Input Jobs

Deploy to ... + **ADD JOB**

Deploy to O

Deploy configura

Deployer type

Select job type

Build

Deploy


Test

Change the name to **Check Deployment Health** and deployment type to **Kubernetes**. Select your API Key from Key Protect the same way as you did above.


Scrolling down, enter US South as the **IBM Cloud Region**. Select the **Resource Group** where your cluster is created which should be **zero-to-cloud-native**. Next, selected your **cluster name** which should be **zero-to-cloud-native**.

For the deploy script, enter:

```
source ./pipeline/scripts/check_health.sh
```


Check Deployment Health  Remove

Deploy configuration


Deployer type  ①

Kubernetes ▼


Pipeline image version ①

Inherited from Configure Pipeline (latest)  ▼

API key ①

API Key for zero to cloud native toolchain 


IBM Cloud region ①

US South  ▼


Account Name ①

kevin collins's Account ▼


Resource Group ①

zero-to-cloud-native  ▼

Cluster name ①


zero-to-cloud-native  × ▼




Deploy script ①

source ./pipeline/scripts/check_health.sh 

After entering all the information, click on the Workers tab.

DEPLOY Delete

Input **Jobs**  Workers New Environment pro...

 >   ADD JOB


Deploy to ... Check Dep... ADD JOB

On the next screen, select the Private Worker that we created.

DEPLOY Delete

Input Jobs **Workers** New Environment pro...

By default, jobs in a pipeline stage are run using a pool of IBM managed shared workers in the region where the pipeline is defined. You can also define private workers running in your own infrastructure and use them to run the jobs in this stage. [Learn more](#).

Jobs in this stage are run using  🔗 ①

(Private worker) dev-frontend-api-02cn-private-worker ▼

Cancel Save

The final step for this job is to set environment properties. Click on the Environment properties tab.

DEPLOY

Input Jobs **Workers** New Environment pro... Delete

By default, jobs in a pipeline stage are run using a pool of IBM managed shared workers in the region where the pipeline is defined. You can also define private workers running in your own infrastructure and use them to run the jobs in this stage. [Learn more](#).

Jobs in this stage are run using 🔗 ⓘ

(Private worker) dev-frontend-api-02cn-private-worker

Cancel Save

There are three Environment Properties we will need to add. To add a property, click on Add Property and enter the property type.

Add property ⊕

- Text property
- Text area property
- Secure property
- Properties file

Referring to the table below, create the following environment properties.


Property Type	Property Name	Value
Text Property	CLUSTER_NAMESPACE	zero-to-cloud-native
Text Property	DEPLOYMENT_FILE	./deployments/deployment.yaml
Properties File	buildprops	build.properties

After entering all of these properties, click Save.

DEPLOY

Input Jobs Workers New **Environment pro...** Delete

Add property ⊕

⋮	CLUSTER_NAMESPACE	zero-to-cloud-native
⋮	DEPLOYMENT_FILE	./deployments/deployment.yaml
⋮	 buildprops	build.properties

Cancel Save

That completes setting up our classic toolchain with a private worker for the api frontend! Now, let's test it.

Toolchains / [dev-classic-api-frontend-02cn](#) / [dev-classic-frontend-api-toolchain](#) /

dev-classic-frontend-api-toolchain | Delivery Pipeline

[Add Stage](#) [Act](#)

BUILD

STAGE NOT RUN

LAST INPUT

Git URL [🔗](#)

Not yet run

JOBS

[View logs and history](#)

🕒 Fetch Code

Not yet run

🕒 Unit Test

Not yet run

LAST EXECUTION RESULT

No results

>

CONTAINERIZE

STAGE NOT RUN

LAST INPUT

Stage: BUILD / Job: Fetch ...

Not yet run

JOBS

[View logs and history](#)

🕒 Check Dockerfile

Not yet run

🕒 Container Registry

Not yet run

🕒 Build Container Image

Not yet run

LAST EXECUTION RESULT

No results

>

DEPLOY

STAGE NOT RUN

LAST INPUT

Stage: CONTAINERIZE / Jo...

Not yet run

JOBS

[View logs and history](#)

🕒 Deploy to OpenShift

Not yet run

🕒 Check Deployment Health

Not yet run

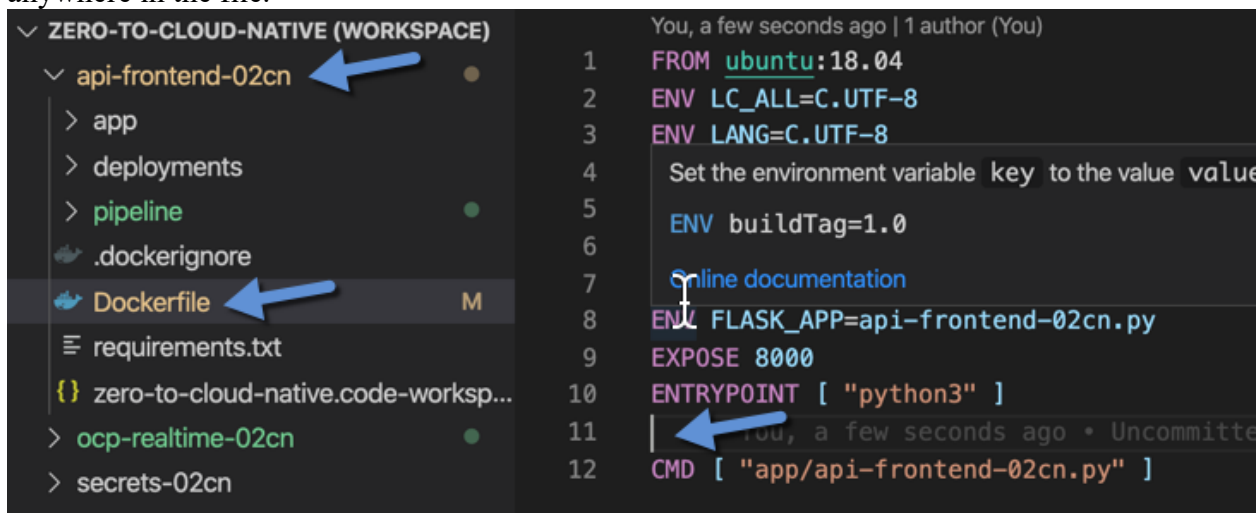
LAST EXECUTION RESULT

No results

4 – 6 Testing the Classic Delivery Pipeline

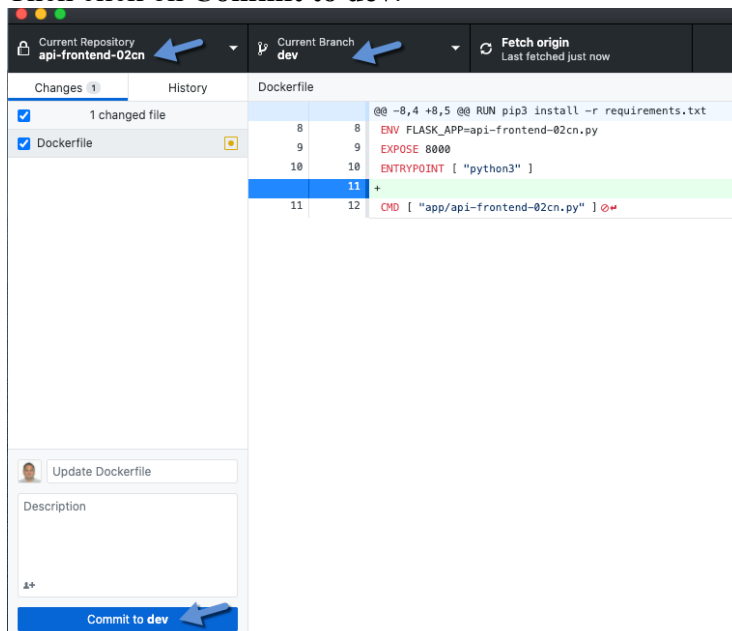
To test our pipeline, all we have to do is make a change to the source code and commit it to GitHub as we have configured our pipeline to automatically start after a GitHub commit to our repository.

To make a change, start Visual Studio Code, and open the workspace you created in step 5. Navigate to the `api-frontend-02cn` folder and click on the `Dockerfile`. Simply add a new line anywhere in the file.

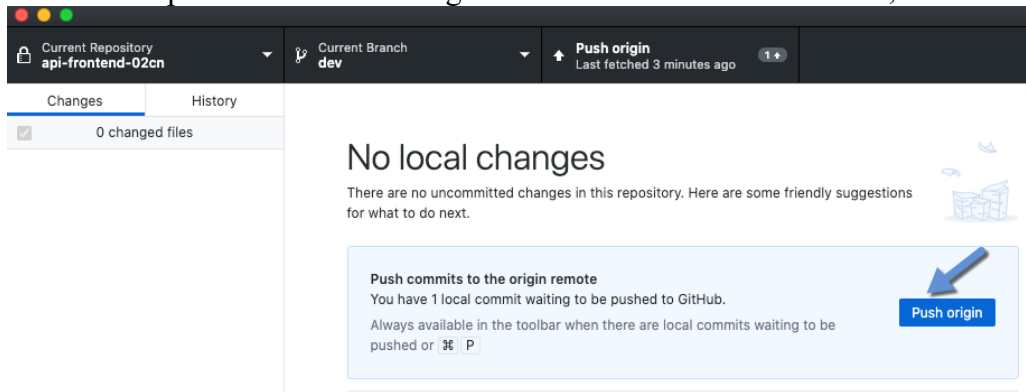


Next, we will need to push the ‘change’ to GitHub. Start your **GitHub Desktop**. Change to the **api-frontend-02cn** as the repository. Make sure the **current branch** is dev.

You will notice the highlighted changes which should only be adding the newline. Then click on **Commit to dev**.



The final step is the Push the changes to GitHub. On the next screen, click on Push origin.



Now, navigate back to your browser with the api frontend pipeline and watch it run!